

# Task 6 - Use orchestration

Hongyi Zhang <hongyiz@kth.se>

Lida Liu <lidal@kth.se>

## I. INTRODUCTION

The Hadoop Distributed File System is designed to be suitable for distributed file systems running on commodity hardware. There are two types of nodes in the HDFS architecture, one is NameNode, also called "metadata node". The other is DataNode, also called "data node". These two types of nodes respectively perform the execution nodes of the specific tasks of Master and Worker. The general design idea: divide and conquer - distribute large batches of files on a large number of independent servers, in order to take the divide and conquer way of computing and analyzing massive data.

## II. IMPLEMENTATION OF HOPSFS

We have encountered several troubles during the implementation. The version of hops-util located in pom.xml in vt18-id2210-cloud-master/hops/ should be 0.1.0, not \${hops.version}. Otherwise the maven will throw the error since it cannot find the corresponding version in SICs Release Repository.

And while we are doing ssh and scp operation, we could also include the private ssh key to accomplish the command.

After all the configuration and installation have been finished, we can get the result as followed.

```
[ubuntu@gceharry-hopsfs-nn-1:~$ /srv/hops/hadoop/bin/hdfs dfs -ls /test
Found 1 items
-rw-r--r-- 1 ubuntu ubuntu 0 2018-05-17 12:30 /test/test.txt
ubuntu@gceharry-hopsfs-nn-1:~$
```

Fig. 1. The file created in HopeFS

Also there are three virtual machines deployed on Google Cloud Platform.

<input type="checkbox"/>	<input checked="" type="checkbox"/>	gceharry-hopsfs-ndbd-1	europe-west1-b	10.132.0.3	35.190.215.29	SSH	⌵	⋮
<input type="checkbox"/>	<input checked="" type="checkbox"/>	gceharry-hopsfs-ndbd-2	europe-west1-b	10.132.0.4	35.205.241.19	SSH	⌵	⋮
<input type="checkbox"/>	<input checked="" type="checkbox"/>	gceharry-hopsfs-nn-1	europe-west1-b	10.132.0.2	35.205.73.148	SSH	⌵	⋮

Fig. 2. Virtual Machines

The following figure is the initial status of three machines on the karamel terminal.

HopsFS is DAG_JOB									
Passed Phases									
Phase	Status	Duration							
1. PRECLEANING	SUCCESS	56							
2. FORMING_GROUPS	SUCCESS	465							
3. FORMING_MACHINES	SUCCESS	1750							
4. FORMING_DAG	SUCCESS	31447							
Tasks' Status									
OS Family	Public IP	Private IP	SSH Port	SSH User	Life Status	Task Status			
UBUNTU	35.190.215.29	10.132.0.3	22	ubuntu	CONNECTED	EMPTY			
Task									
Task	Status	Actions	Duration(ms)						
1. find os-type	DONE	log	482						
2. apt-get essential	DONE	log	17777						
3. install chdfs	DONE	log	13334						
4. make solo-zb	DONE	log	141						
5. clone and vendor hops-hadoop-chef	DONE	log	21932						
6. hops-install	DONE	log	46188						
7. hops-ids	DONE	log	23879						
8. ndbd-install	DONE	log	47624						
9. ndbd-ndbd	DONE	log	6559						

OS Family									
OS Family	Public IP	Private IP	SSH Port	SSH User	Life Status	Task Status			
UBUNTU	35.205.241.19	10.132.0.4	22	ubuntu	CONNECTED	EMPTY			
Task									
Task	Status	Actions	Duration(ms)						
1. find os-type	DONE	log	482						
2. apt-get essential	DONE	log	17777						
3. install chdfs	DONE	log	13334						
4. make solo-zb	DONE	log	141						
5. clone and vendor hops-hadoop-chef	DONE	log	21932						
6. hops-install	DONE	log	46188						
7. hops-ids	DONE	log	23879						
8. ndbd-install	DONE	log	47624						
9. ndbd-ndbd	DONE	log	6559						

Fig. 3. Status of Virtual Machines

## III. BENCHMARK OF HOPSFS

In this section we will benchmark the throughput of the HopeFS System. We will use 10 files with 557.47mb in total. With command of the File System (FS) shell, we can easily benchmark the reading and writing ability for the HDFS. The code for this section can also be found on Github.

The following figure shows all the files which has been stored in the HopeFS.

```
[ubuntu@gceharry-hopsfs-nn-1:~$ /srv/hops/hadoop/bin/hdfs dfs -ls /rwtest
Found 10 items
-rw-r--r-- 3 ubuntu ubuntu 110166656 2018-05-17 13:59 /rwtest/Advanced_Internetworking.zip
-rw-r--r-- 3 ubuntu ubuntu 5108718 2018-05-17 13:59 /rwtest/102218-Blockchains.pdf
-rw-r--r-- 3 ubuntu ubuntu 166839811 2018-05-17 13:59 /rwtest/102218.zip
-rw-r--r-- 3 ubuntu ubuntu 29334849 2018-05-17 13:59 /rwtest/International_Consulting.zip
-rw-r--r-- 3 ubuntu ubuntu 99783428 2018-05-17 13:59 /rwtest/Internet_Security.pdf
-rw-r--r-- 3 ubuntu ubuntu 88566112 2018-05-17 13:59 /rwtest/JAVA.zip
-rw-r--r-- 3 ubuntu ubuntu 88482229 2018-05-17 13:59 /rwtest/Java_Programming.pdf
-rw-r--r-- 3 ubuntu ubuntu 8126464 2018-05-17 13:59 /rwtest/Line.sdf
-rw-r--r-- 3 ubuntu ubuntu 9156764 2018-05-17 13:59 /rwtest/Project.zip
-rw-r--r-- 3 ubuntu ubuntu 167177782 2018-05-17 13:59 /rwtest/Research_Methodology_and_Scientific_Writin
g.zip
```

Fig. 4. Files in HopeFS

Here is the result for reading test. For reading operation, we are using FS shell copyToLocal command `hdfs dfs -copyToLocal [-ignorecrc] [-crc] URI <localdst>`. At first we want to use the moveToLocal but this command haven't been implemented yet.




Fig. 5. Reading Throughput

The throughput is pretty high. The average reading rate is 21.76mb/s while the maximum can reach up to 49.51mb/s. The maximum rate is the transmission for the file *Research\_Methodology\_and\_Scientific\_Writing.zip*, which is also the maximum-size file. That means the HopeFS is really suitable to handle large files. You can also see the result in writing operation.




Fig. 6. Writing Throughput

The writing rate is really similar to reading process. The average is 25.23mb/s while the maximum can be 62.65mb/s also for that largest file. Besides that, all the large size file have a pretty high speed rate.

As HDFS is a Master/Slave architecture. From the end user's point of view, it can perform CRUD (Create, Read, Update, and Delete) operations on a file through the directory path just like a traditional file system. However, due to the nature of distributed storage, HDFS clusters have a NameNode and some DataNodes. The NameNode manages the metadata of the file system and the DataNode stores the actual data. The client accesses the file system through interactions with the NameNode and DataNodes. The client contacts the NameNode to get the metadata of the file, and the real file I/O operation is directly interacting with the DataNode.

Since HDFS is generally used for "write once, read many times". We think it is not suitable for real-time interaction, and is not suitable for storing a large number of small files.

## REFERENCES

- [1] Apache Hadoop [Online]. Available: <http://hadoop.apache.org>
- [2] Apache Hadoop Project Dist POM, Apache Hadoop 2.4.1 [Online]. Available: <https://hadoop.apache.org/docs/r2.4.1/hadoop-project-dist/hadoop-common/FileSystemShell.html>